CLEF HIPE-2026: Participation Guidelines

Shared Task on Person-Place Relation Extraction from Historical Documents

v.2025-12-04

HIPE-2026 Team:
Juri Opitz, Maud Ehrmann, Simon Clematide,
Corina Raclé, Emanuela Boros, Andrianos Michail, Matteo Romanello

Contents

1 Person–Place Relation Extraction from Multilingual Historical Texts **2**

2 Task Description **2**

3 Data **5**

4 Evaluation Campaign and System Responses **8**

5 Submission Instructions **9**

6 Evaluation **10**

7 CLEF Workshop and Notebook Papers **13**

A Data Preparation **15**

Useful Links

- HIPE-2026 website: `https://hipe-eval.github.io/HIPE-2026/`
- Google Group: `https://groups.google.com/g/hipe-2026`
- HIPE-2026 data: `https://github.com/hipe-eval/hipe-2026-data`
- HIPE-2026 evaluation: `https://github.com/hipe-eval/hipe-2026-eval`

1 Person–Place Relation Extraction from Multilingual Historical Texts

HIPE-2026 is a CLEF Evaluation Lab dedicated to the extraction of person–place relations in multilingual historical documents. A person-place relation corresponds to a semantic link between an individual and a location as evidenced in a document. Such relations may indicate where a person is said to be at a given moment, where they lived or worked, or places connected to notable moments in their life (e.g., birthplaces, residences, visits, travel destinations). Together, these relations can support the reconstruction of individuals' geographical and temporal trajectories.

Detecting these implicit or explicit, spatio-temporal relations cannot be achieved through simple document co-occurrence of entity mentions. Rather, it requires temporal reasoning, geographical inference, and interpretation of noisy historical texts–often with sparse or indirect contextual cues–to detect and qualify person–place relations with appropriate degrees of certainty.

The objective of HIPE-2026 is to advance the automatic detection of such relations, enabling the reconstruction of individuals' movements in space and time and the tracing of life trajectories in support of digital humanities scholarship. The task is designed to be approachable by both generative AI systems (LLMs) and more traditional classification models.

2 Task Description

In essence, participants are tasked with developing systems that determine the nature of the relationship implied by each person-place pair in a historical document. Each pair consists of two entities, one person and one location, each of which appears in the text through one or more mentions. Specifically, systems must establish whether the text indicates that the person is currently at that place within the document's temporal horizon (a recency-oriented **`isAt`** relation), suggests that the person was at that place at some point in time (a more general **`at`** relation), or provides no meaningful evidence linking the person and the place. As illustrated in Figure 1, the `isAt` relation is bounded in time and close to the document publication date, while the `at` relation is only right-bounded by the publication date: the relation could hold for any time in the past, but naturally not in the future. The next sections provide a more detailed and formal definition of the task.

2.1 Terminology

Before defining the sub-tasks, we introduce the terminology. Each term corresponds to fields in the data structure presented in Section 3.2.

- ***context***: The full text of a document or passage together with its metadata (for example, newspaper title, language, publication date).

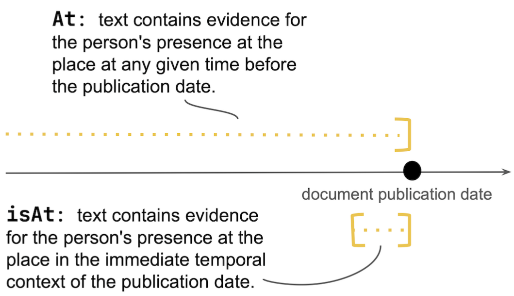

Figure 1: Schematic representation of the two person-place relation types.

- ***person***: An entity, represented by a cluster of mentions in a document, referring to a specific person.
- ***location***: An entity, represented by a cluster of mentions in a document, referring to a specific location. We use the terms ‘location‘ and ‘place‘ interchangeably.
- ***(person, location)***: A pair consisting of one person entity and one location entity drawn from the same document. Not all possible pairs are included in order to avoid excessively large candidate sets.
- **`at`**: A relation type indicating that the text provides evidence that the person was at the location at some point in time.
- **`isAt`**: A temporally narrower relation type indicating that the text suggests the person is at the location within the temporal horizon of the news document (e.g., the time of reporting).
- ***system***: A system developed by participants that, given a context, a $(person, location)$ pair, and a candidate relation type, assigns a label to that candidate relation.

2.2 Sub-tasks

The shared task comprises two sub-tasks, each defined by one of the two relation types that systems must classify.

Sub-task definition: `at` relation

Given a *context* and a $(person - location)$ pair, the general **`at`** relation captures whether the context implies that the person was at that place at any point in time. Systems should perform the following classification:

$$\text{system}\left[\mathit{context}, (\mathit{person}, \mathit{location}), \texttt{at}\right] = \begin{cases} \textsc{true} \text{ if the context provides explicit evidence that the person was at the place at some point in time.} \\ \textsc{probable} \text{ if this can be \textit{inferred} from \textit{implicit} cues in the context and is thus a likely assumption.} \\ \textsc{false} \text{ if no evidence is present or the context contradicts such relation.} \end{cases}$$

Sub-task definition: `isAt` relation

Given a *context* and a $(\mathit{person}, \mathit{location})$ pair, the temporally narrower **`isAt`** relation captures whether the context implies that the person was at that place within the document's temporal horizon. Systems will perform the following classification:

$$\text{system}\left[\mathit{context}, (\mathit{person}, \mathit{location}), \texttt{isAt}\right] = \begin{cases} \textsc{true} \text{ if the person is located at the place shortly before the document's publication time.} \\ \textsc{false} \text{ if this is not the case.} \end{cases}$$

Temporal window for `isAt`. By definition, `isAt` = TRUE means there is evidence the person was at the location up to about one month before the publication date. Outside this window, it should be set to FALSE.

Relation between `at` and `isAt`. The `isAt` relation is a temporal refinement of `at`: it captures whether the `at` relation holds shortly before or at the time of publication. The following constraints apply. If `at` is FALSE, then assigning TRUE to `isAt` would be inconsistent. If `at` is TRUE or PROBABLE, `isAt` further specifies whether the person's presence falls within the document's temporal horizon.

For example, if an article reports a current visit of a US President to Berlin, then the relation between the president and Berlin is labeled as `at` and `isAt`. By contrast, if the article refers to a past trip ("Last year, the president visited Berlin."), only `at` applies, since the event lies outside the temporal horizon of the document.

For simplicity, however, the shared task evaluates the two relations independently.

3 Data

3.1 Data Overview

HIPE-2026 data consists of two sets of data: historical newspaper articles in French, German, and English spanning roughly 200 years (19C-20C), and literary documents in French. The newspaper material–originating from European and US library collections–composes the core dataset of the shared task and is released as training, development, and test sets ('Test A'). The literary documents are reserved as a surprise test set ('Test B') to evaluate systems' robustness and generalization on out-of-domain data.

The historical newspaper data is derived from the entity-annotated HIPE-2022 material and includes only those datasets that contain *person* and *location* annotations. Data has been prepared and newly annotated for the present shared task. The literary surprise set has likewise been newly curated. Appendix 7 details the preparation of the historical newspaper data, while a description of the surprise set preparation will be released after the evaluation campaign.

3.2 Data Representation and Contents

Data are represented according to a dedicated JSON schema[1] and stored as JSON Lines (.jsonl) files, where each line corresponds to one document. Each file contains a collection of documents belonging to a single language (German, French, or English). These documents correspond to individual historical newspaper articles or literary text chunks, represented independently. Every document consists of four main components presented below and illustrated in Figure 2:

- **Metadata keys** provide contextual information about the document and includes:
  - `document_id`: a unique identifier for the document;
  - `media`, including:
    - `publication_title`;
    - `time_period`: the year interval spanned by the media sample in the dataset;
    - `source_type`: e.g. "newspaper";
  - `source`: data path in the original material;
  - `language`: the language of the document (one of de, fr, or en);
  - `date`: the article's publication date.
- The **`text`** key contains the full text of the document. This is the context within which persons and locations are mentioned, and from which systems must infer whether person–place relations hold.
- The **`sampled_pairs`** list contains person–location pairs selected for annotation. Each pair entry includes:

[1] https://github.com/hipe-eval/HIPE-2026-data/blob/main/schemas/hipe-2026-data.schema.json

```
{ "document_id": "NZZ-1798-12-12-a-p0003",
  "media": {
    "publication_title": "Neue Zürcher Zeitung",
    "time_period": "1780-1950",
    "source_type": "Newspaper"
  },
  "source": "v2.1/hipe2020/de/HIPE-2022-v2.1-hipe2020-dev-de.tsv",
  "language": "de",
  "date": "1798-12-12",
  "text": "(...) Der be\nkannte Irländer Theobald Wolfe Tone, den man auf
  ↪  der\nBompartschen Eskader gefangen nahm, ward von einem\nKriegsgericht zum Tode
  ↪  verurtheilt. Er hörte sein Ur\ntheil mit Standhaftigkeit an , und verlangte nur,
  ↪  daß\nman ihn, seiner franz. Uniform halber, erschießen soll\nte. Da ihm das
  ↪  versagt ward, verwundete er sich mit\neinem Messer so, daß er schwerlich aufkommen
  ↪  wird. -\nDie von Nelson bey Abukir Lissabon angekommen; nur das Schif (...)",
  "sampled_pairs": [
    {
      "pers_entity_id": "NZZ-1798-12-12-a-p0003_Q83235",
      "pers_wikidata_QID": "Q83235",
      "pers_mentions_list": ["Nelsons", "Nelson", "Adm. Nelson", "Baron Nelson vom\nNil
      ↪  und Burnharm Torpe"
      ],
      "loc_entity_id": "NZZ-1798-12-12-a-p0003_Q12783927",
      "loc_wikidata_QID": "Q12783927",
      "loc_mentions_list": ["Abukir"],
      "at": "TRUE",
      "at_explanation": "(...)",
      "isAt": "FALSE",
      "isAt_explanation": "(...)"
    },
    {
      "pers_entity_id": "NZZ-1798-12-12-a-p0003_Q335205",
      "pers_wikidata_QID": "Q335205",
      (...)
```


Figure 2: **Example of a HIPE-2026 input document.** Key values are shortened for readability. In the training and development sets, the `at/isAt` fields contain gold labels, whereas in the test set they are set to `null`. Systems must predict the values for `at` and `isAt`, and may optionally provide explanations via the `*_explanation` fields.

Data Source	Lang.	#Docs	#Pairs
`hipe2020`	de	146	1,974
	fr	214	3,055
	en	94	827
`sonar`	de	13	180
`newseye`	de	4	64
	fr	77	1,097
`letemps`	fr	256	3,470
Total	–	804	10,667

Split	Lang.	#Docs	#Pairs
train	de	97	1,345
	fr	328	4,521
	en	56	501
dev	de	32	432
	fr	109	1,557
	en	18	148
test	de	34	441
	fr	110	1,544
	en	20	178
Total		804	10,667

Table 1: HIPE-2026 historical newspaper dataset statistics per data source (left) and per language split (right). Values reflect pre–manual-annotation counts and represent **upper bounds**; final released data will contain less pairs.

- `pers_entity_id` and `loc_entity_id`: unique identifiers of the entities;
- `pers_wikidata_QID` and `loc_wikidata_QID`: Wikidata identifiers, when available;
- `pers_mentions_list` and `loc_mentions_list`: lists of mentions for the entities;
- `at` and `isAt`: relation labels (gold value in training/dev; null in test/surprise);
- `at_explanation` and `isAt_explanation`: optional explanations for the assigned labels.

3.3 Data Splits and Statistics

For the historical newspaper data, training, development, and test sets are provided. For the literary data, only the surprise test set is released. Table 1 reports statistics for the prepared data (prior to full manual annotation at the time of publication), and therefore reflects the upper bound on the number of person–place pairs that will be available. This section will be updated in next version of the guidelines.

3.4 Data releases

The HIPE-2026-data directory is organized per dataset, release version and language, as shown in Figure 3.

- There is one UTF-8 encoded, .jsonl file per language and split.
- All files, including system submissions, adhere to the same JSON schema.
- Files are named according to this schema:
  `HIPE-2026-<hipeversion>-impresso-<split>-<language>.jsonl`, where the value of split can be `sample`, `train`, `dev`, or `test`.

- HIPE-2022 releases are versioned with a two-part version number (Major.Minor) present in 1) the data directory structure and 2) the filename of each file.
- Each HIPE-2026 release has an equivalent git repository release, with release notes.

```
data
└── newspapers
    ├──  vx.x
         ├── lg1
         |      ├── HIPE-2026-vx.x-newspapers-train-lg1.jsonl
         |      ├── HIPE-2026-vx.x-newspapers-dev-lg1.jsonl
         └── lg2
         |      ├── HIPE-2026-vx.x-newspapers-train-lg2.jsonl
         |      ├── HIPE-2026-vx.x-newspapers-dev-lg2.jsonl
         ├── ...
    └── vx.x
         ├── lg1 ...
  └── literaryworks
      ├──  vx.x
             ├── lg1
             |      ├── HIPE-2026-vx.x-literaryworks-test-lg1.jsonl
```

Figure 3: HIPE-2026 release data directory structure.

4 Evaluation Campaign and System Responses

4.1 Evaluation Campaign General Rules

- Registration is open until 23 April 2026. Please refer to the HIPE timeline for more information.

- The evaluation period runs from Tue 5 May 2026, 10 a.m. (CEST) to Thu 7 May 2026 8 p.m (CEST). At the beginning of this period, input test files for all languages and the surprise test set will be released on the hipe-data-2026 repository. At the end of the evaluation period, participants must send their system responses via email to the submission email (see below). Gold-standard data will be released after the publication of the evaluation results.

- Teams may submit up to **three** runs per language.

- Teams must provide predictions for all input test files (Test A for all languages + surprise Test B).

- Teams may use any external resources (e.g. additional language resources available elsewhere or homemade, and other annotated data).

5 Submission Instructions

Participants are required to submit system responses in JSONL format, conforming to the HIPE-2026 schema. Each response file corresponds to a system run and must include predictions for all documents of a given input test set.

5.1 Submission Task Overview

Each test input file contains a list of documents, each with sampled person–location pairs. For each such pair, systems must:

- Replace `null` values in the `at` field with one of: `"TRUE"`, `"PROBABLE"`, or `"FALSE"`.
- Replace `null` values in the `isAt` field with either: `"TRUE"` or `"FALSE"`.
- Optionally, provide a textual explanation by filling the corresponding `at_explanation` and `isAt_explanation` fields.

Important: Any missing documents, or `null` label left in the `at` or `isAt` fields at submission time will be interpreted as `"FALSE"` during evaluation.

5.2 Task Scope

- Both sub-tasks (`at`, `isAt`) are evaluated on the historical newspaper Test A.
- Only the `at` sub-task is evaluated on the surprise Test B (literary domain).
- Systems must return predictions for all languages of Test A. The surprise test set B is evaluated separately.
- Teams may submit up to **three** runs per test set (maximum of 6 runs in total for both test sets A and B).

5.3 File Format and Schema Compliance

Each system run must be submitted as a single UTF-8 encoded `.jsonl` file that:

- Mirrors the structure and document ordering of the input test file.
- Preserves all document and entity metadata.
- Adheres strictly to the HIPE-2026 data schema hipe-2026-data.schema.json

5.4 File Naming and Packaging

- Response files must be named using the following pattern:

  ```
  TEAMNAME_<inputfile>_runX.jsonl
  ```

  where `TEAMNAME` is your registered CLEF team name, and `runX` is one of `run1`, `run2`, or `run3`.

- All run files must be packaged into a single ZIP archive named:

  ```
  TEAMNAME.zip
  ```

- Submit the ZIP archive via email to: `hipe-2026-submission@listes.epfl.ch`
- You will receive a confirmation of receipt.

5.5 Validator and Support Tools

Participants are encouraged to use the provided:

- JSON Schema validator to check formatting and schema compliance.
- Scorer tool to compute macro Recall and verify system behavior.

Both tools are available in the hipe-2026-eval repository.

6 Evaluation

Three evaluation profiles are considered:

1. The **accuracy profile** rewards high-performance systems, encouraging exploration of frontier models, advanced prompt engineering, or agentive approaches.
2. The **efficiency profile** promotes lightweight, scalable methods, including smaller LLMs or task-specific classifiers.
3. The **generalization profile** tests system accuracy on a surprise test set in a non-news domain, based on the `at`-relation only.

6.1 Accuracy Evaluation Profile

Metric

We evaluate system responses per test set using *macro Recall* (also known as *balanced accuracy*). Macro Recall is defined as follows:

1. For each label ℓ in the label set L, compute its recall:

$$\text{Recall}(\ell) = \frac{\#\text{examples with label } \ell \text{ correctly predicted}}{\#\text{examples whose gold label is } \ell}.$$

2. The final score is the arithmetic mean of the per-label recalls:

$$\text{MacroRecall} = \frac{1}{|L|}\sum_{\ell \in L} \text{Recall}(\ell).$$

This metric is highly interpretable and ensures that all labels contribute equally, regardless of class imbalance [8, 7].

Interpretation of `null` value An `at` or `isAt` label value of `null` means "system makes no assertion." During evaluation, any `null` is converted to FALSE for both `at` and `isAt`.

Evaluation of Test A

The Main Test A evaluates each sub-task independently, both scored with macro Recall:

- `at`: A 3-class classification problem, evaluated using macro Recall ($|L| = 3$).
- `isAt`: A 2-class classification problem, evaluated using macro Recall ($|L| = 2$).

The global test score for Test A is computed as follows:

$$\text{GlobalScore}_A = \frac{\text{MacroRecall}_{\text{At}} + \text{MacroRecall}_{\text{isAt}}}{2}.$$

The scorer is applied *separately* to:

1. the predictions for the `at` sub-task,
2. the predictions for the `isAt` sub-task.

It then automatically computes and reports:

- the macro Recall for each sub-task, and
- the global average score defined above.

6.2 Accuracy-Efficiency Evaluation Profile with Main Test Set B

HIPE-2026 features a pilot evaluation that balances system accuracy with efficiency and resource usage. This is motivated by the increasing cost of running very large models and, in our case, by the scale of digitized material together with the quadratic nature of the task (person–location *pairs*). We aim to understand how efficient and scalable approaches can be designed. To this end, we will survey participant teams for the following two types of efficiency information:

1. **Parameter Count** Params(s): the total number of model parameters.
2. **Model Size** Size(s): the disk footprint of the deployed model (in MB or GB).

These dimensions capture two aspects of model efficiency: conceptual complexity (parameter count) and storage or deployment cost (model size). Both metrics must be reported by participants following the provided guidelines.

Scoring and Ranking

We design a robust accuracy-efficiency metric. Let $\mathcal{S}$ be the set of submitted systems. For each dimension $d \in \{\text{Params}, \text{Size}, \text{Acc}\}$, where Acc denotes the system's official accuracy metric (macro Recall), each system receives a rank $r_d(s)$:

$$r_d(s) = \text{rank of system } s \text{ on dimension } d,$$

with lower ranks indicating better performance (1 = best). For efficiency dimensions, smaller values are strictly preferred (fewer parameters, smaller model size).

The final composite score for system s is the arithmetic mean of its ranks across the three dimensions:

$$R(s) = \frac{1}{3}\left(r_{\text{Acc}}(s) + r_{\text{Param}}(s) + r_{\text{Size}}(s)\right).$$

The systems are globally ordered by increasing $R(s)$.

Worked Example

Consider three hypothetical systems with the following properties:

System	Accuracy	Parameter Count	Model Size
s_2	0.79	350M	1.2 GB
s_3	0.77	1.3B	5.0 GB

Step 1: Ranking each dimension.

$$\begin{aligned}
\text{Accuracy ranks:} \quad & r_{\text{Acc}}(s_1) = 1,\ r_{\text{Acc}}(s_2) = 2,\ r_{\text{Acc}}(s_3) = 3. \\
\text{Param. count ranks (fewer is better):} \quad & r_{\text{Param}}(s_2) = 1,\ r_{\text{Param}}(s_3) = 2,\ r_{\text{Param}}(s_1) = 3. \\
\text{Model size ranks (smaller is better):} \quad & r_{\text{Size}}(s_2) = 1,\ r_{\text{Size}}(s_3) = 2,\ r_{\text{Size}}(s_1) = 3.
\end{aligned}$$

Step 2: Final aggregated score.

$$\begin{aligned}
R(s_1) &= \frac{1+3+3}{3} = 2.33, \\
R(s_2) &= \frac{2+1+1}{3} = 1.33, \\
R(s_3) &= \frac{3+2+2}{3} = 2.33.
\end{aligned}$$

Final ordering.

$$s_2 \prec \{s_1, s_3\}$$

System s_2 ranks highest overall due to its substantially smaller parameter count and compact model size, despite not achieving the highest accuracy.

6.3 Generalization Profile: Surprise Test B

The Surprise Test B assesses the ability of systems to generalize beyond the historical newspaper domain. It consists of literary documents that differ in style, genre, and temporal expression, posing additional challenges for relation extraction.

This evaluation profile includes only the `at` relation, framed as a 3-class classification task {TRUE, PROBABLE, FALSE}.

Evaluation is conducted using **macro Recall**, ensuring equal weight across all three classes regardless of class distribution.

A single scorer run is required. It outputs the following metric:

$$\mathrm{MacroRecall}_{\texttt{at}}$$

This generalization profile allows for the analysis of system robustness and adaptability to unseen, non-news textual data.

6.4 Optional LLM Explanations

Explanations (`*_explanation`) are ignored during scoring and have no influence on any metric. Participants may use them to understand LLM-based approaches rationale and discuss them. We may use them to qualitatively analyze and inspect results and present possible findings in the final overview paper.

6.5 Technical Scoring Procedure

We provide an output validator and a scorer in our GitHub repository. The output validator checks if the submitted file has the right content and schema, validating it against our JSON schema. The scorer scores the above measures, and provides an evaluation report that helps the participants to track system progress.

7 CLEF Workshop and Notebook Papers

Participants will submit a Notebook Paper to be presented during the final HIPE-2026 workshop co-located with the CLEF conference (September 2026) and to be published online via the CEUR Workshop Proceeding open access publication service. Please check submission instructions and important dates on the HIPE website. Previous HIPE-eval proceedings are available: see HIPE-2022 and HIPE-2020.

Acknowledgments and Organisation

The HIPE-2026 organising team expresses its sincere appreciation to the CLEF-2026 Evaluation Lab Organising Committee for the overall coordination and support. HIPE-eval editions are organised within the framework of the Impresso - Media Monitoring of the Past project, funded by the Swiss National Science Foundation under grant No. CRSII5_213585 and by the Luxembourg National Research Fund under grant No. 17498891.

HIPE-2026 is organised by Juri Opitz (UZH), Corina Raclé (UZH), Emanuela Boros (EPFL), Andrianos Michail (UZH), Matteo Romanello (UZH), Maud Ehrmann (EPFL), Simon Clematide (UZH).

References

[1] M. C. Ardanuy, D. Beavan, K. Beelen, K. Hosseini, J. Lawrence, K. McDonough, F. Nanni, D. van Strien, and D. C. S. Wilson. A Dataset for Toponym Resolution in Nineteenth-Century English Newspapers. *Journal of Open Humanities Data*, 8(0):3, Jan. 2022.

[2] M. Ehrmann, M. Romanello, A. Flückiger, and S. Clematide. Impresso Named Entity Annotation Guidelines. Annotation guidelines, Ecole Polytechnique Fédérale de Lausanne (EPFL) and Zurich University (UZH), Jan. 2020.

[3] M. Ehrmann, M. Romanello, S. Najem-Meyer, A. Doucet, and S. Clematide. Overview of HIPE-2022: Named Entity Recognition and Linking in Multilingual Historical Documents. In *Experimental IR Meets Multilinguality, Multimodality, and Interaction: 13th International Conference of the CLEF Association, CLEF 2022, Bologna, Italy, September 5–8, 2022, Proceedings*, pages 423–446, Berlin, Heidelberg, Sept. 2022. Springer-Verlag.

[4] A. Hamdi, E. Linhares Pontes, E. Boros, T. T. H. Nguyen, G. Hackl, J. G. Moreno, and A. Doucet. A Multilingual Dataset for Named Entity Recognition, Entity Linking and Stance Detection in Historical Newspapers. In *Proceedings of the 44th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 2328–2334, New York, NY, USA, July 2021. Association for Computing Machinery.

[5] S. Menzel, H. Schnaitter, J. Zinck, V. Petras, C. Neudecker, K. Labusch, E. Leitner, and G. Rehm. Named Entity Linking mit Wikidata und GND – Das Potenzial handkuratierter und strukturierter Datenquellen für die semantische Anreicherung von Volltexten. In *Named Entity Linking mit Wikidata und GND – Das Potenzial handkuratierter und strukturierter Datenquellen für die semantische Anreicherung von Volltexten*, pages 229–258. De Gruyter Saur, Oct. 2021.

[6] S. Menzel, J. Zinck, H. Schnaitter, and V. Petras. Guidelines for Full Text Annotations in the SoNAR (IDH) Corpus. July 2021.

[7] J. Opitz. A closer look at classification evaluation metrics and a critical reflection of common evaluation practice. *Transactions of the Association for Computational Linguistics*, 12:820–836, 06 2024.

[8] F. Sebastiani. An Axiomatically Derived Measure for the Evaluation of Classification Algorithms. In *Proceedings of the 2015 international conference on the theory of information retrieval*, pages 11–20, 2015.

A Data Preparation

A.1 HIPE-2022 historical newspaper data

Historical newspaper data is derived from the entity-annotated HIPE-2022 material[2] and includes only those datasets that contain *person* and *location* annotations, that is to say: `hipe2020`, `newseye`, `sonar`, and `letemps`.

These datasets consists of historical newspaper material in French, German and English annotated with named entity mentions, each linked to a Wikidata unique identifier (QID) or a NIL entry. For more information on the HIPE-2022 dataset, please refer to [3] for the overview, to [5, 6, 4, 1] for the primary datasets, and to the Impresso named entity annotation guidelines [2].

Starting from HIPE-2022, the preparation of the HIPE-2026 data involves extracting and annotating person–location pairs, and restructuring information in a format better suited for large language model-assisted approaches. The main preparation steps were as follows:

1. **File management and representation transformation**: The original HIPE-2022 data splits are ignored, and all documents are regrouped per language, regardless of their dataset of origin. The tabular separated source files (`.tsv`), which encode named entity information using the IOB tagging scheme, are converted into JSON representations following a JSON schema and stored as JSON line files. Only the relevant information is retained: document metadata, full text, and all entity mentions of type *person* and *location* with their offsets, surface forms, and Wikidata QID or NIL value.

2. **Grouping of mentions into entities**: Mentions sharing the same QID are grouped into a single entity object. NIL mentions are grouped either with other NIL mentions or merged into non-NIL entities when their surface form exceeds a 75% Levenshtein similarity threshold. After this stage, every mention is assigned to one entity object, that has a unique ID (either Wikidata QID or a newly made one for merged NIL). Sanity checks and statistics are computed, and overly long documents are filtered out.

3. **Extraction of candidate person–location pairs**: Person–location pairs are generated at the entity level (each entity object already aggregates all its mentions). For every document, the cross-product of person and location is computed. From this set, a subset of pairs is sampled using a proximity-based score: the minimum number of tokens between any person mention and any location mention. This ensures that most selected pairs are in close textual proximity, while still retaining a few long-distance pairs. The resulting data conforms to the final JSON schema.

4. **Person–location relation annotation**: Candidate pairs are pre-annotated using an LLM and then manually reviewed and corrected collaboratively.

[2] In its version v2.1 available here: `https://github.com/hipe-eval/HIPE-2022-data/tree/main/data/v2.1`

5. **Final dataset creation**: Dataset splits are assembled and packaged for release using the final JSON schema.

A.2 Literary data

Information will be published after the evaluation campaign.